

# Real-time Speech Transcription POC – Working Document

## Overview

- Browser microphone → Client (HTML/JS) → Backend WebSocket (`/ws`) → AWS Transcribe Streaming → live transcript → S3 on stop.
- On Stop: concatenates only final transcripts; uploads `YYYYMMDD\_HHMMSS.txt` to S3 bucket.

## Services and Ports

- Backend: FastAPI at `http://localhost:8000` (WebSocket `/ws`, health `/health`)
- Client: Static HTML at `http://localhost:8080/simple_client.html`

## WebSocket Protocol

- Client → Backend
- `{ "type": "audio", "payload": "" }`
- `{ "type": "control", "action": "stop" }`
- Backend → Client
- `{ "type": "transcript", "text": "...", "is\_final": true|false }`
- `{ "type": "saved", "s3\_key": "YYYYMMDD\_HHMMSS.txt" }`

## Environment Variables

- `AWS\_ACCESS\_KEY\_ID`, `AWS\_SECRET\_ACCESS\_KEY`, `AWS\_REGION`, `S3\_BUCKET`
- Optional: `TRANSCRIBE\_LANGUAGE\_CODE` (default `en-US`), `LOG\_LEVEL` (default `INFO`).

## Docker Compose

- `backend` (port 8000) and `client` (port 8080) services; shared bridge network; backend has healthcheck; client waits for backend healthy.

## Backend (FastAPI + AWS Transcribe)

- File: `backend/main.py`
- Key parts:
- Transcribe Streaming: `start_stream_transcription(language_code=..., media_sample_rate_hz=16000, media_encoding="pcm")``
- WebSocket `/ws`: accepts JSON messages; decodes base64 audio; feeds to AWS stream; on stop, ends stream, concatenates finals, uploads to S3.

- S3 save via ``boto3.put_object`` offloaded with ``run_in_executor``.
- Sends back live/final transcripts and final ``saved`` notification.

## Client (HTML/JavaScript)

- Files: ``client/simple_client.html``, ``client/serve_client.py``, ``client/Dockerfile``
- Flow:
- Opens WebSocket to ``ws://localhost:8000/ws``.
- Captures mic with Web Audio API at 16kHz, mono.
- Converts Float32 to PCM16 and base64; sends as ``{type:"audio"}``.
- Displays live text and accumulates final text.
- On Stop: sends ``{type:"control", action:"stop"}``; shows S3 key on save.

## Terraform

- File: ``terraform/main.tf``
- Provisions:
- S3 bucket (versioned, encrypted, public-blocked).
- IAM user with inline policy: ``transcribe:StartStreamTranscription``, ``s3:PutObject``, ``s3:AbortMultipartUpload`` for that bucket.
- Access keys outputs (sensitive).

## Runbook

1. Terraform: ``cd terraform && cp terraform.tfvars.example terraform.tfvars && terraform init && terraform apply``
2. ``.env``: copy ``env.example`` → ``.env`` and fill outputs (keys, region, bucket)
3. Start: ``docker compose up -d --build``
4. Use: open ``http://localhost:8080/simple_client.html``, Start/Stop to see/save transcripts

## Troubleshooting

- WebSocket fails: ensure backend healthy and ports free; env vars set.
- Mic issues: allow permissions; HTTPS required in production.
- S3 failures: check bucket name and IAM policy.
- CORS: backend allows ``http://localhost:8501`` by default; add ``http://localhost:8080`` if needed.

## Production Notes

- Use HTTPS for client; tighten CORS; prefer IAM roles; add monitoring/logging; scale via ECS/EKS and ALB.